

Avijit Ghosh

Lead Technical AI Policy Researcher at Hugging Face

Research Affiliate at the University of Connecticut

☎ (+1) 857-337-0180 | ✉ avijit@huggingface.co (Work) | avijitg22@gmail.com (Personal) | 🌐 evijit.io | in evijit | 📺 evijit | 📺 evijit

Responsible AI Machine Learning AI Policy AI for Science AI Governance

EDUCATION

- 2019 – 23

Northeastern University
Ph.D. in Computer Science (Advised by Dr. Christo Wilson).
Thesis: *Algorithmic Fairness in the Real World: Challenges and Considerations.*

Boston, MA
- 2014 – 19

Indian Institute of Technology (IIT) Kharagpur
B.Tech. in Chemical Engineering, M.Tech in Financial Engineering, Minor in Computer Science.

Kharagpur, IN

WORK EXPERIENCE

Current Appointments

- 2024 – Now

Hugging Face **Industry**
Lead Technical AI Policy Researcher.
Promoted: Applied → Technical → Lead. Lead HF's external technical-policy engagement, briefing US, UK, EU, and Singapore legislators and regulators on open-source AI governance, evaluation standards, and third-party flaw disclosure.
 - EvalEval Coalition** 2024 – Now
Founder & Lead. Multi-stakeholder coalition of 500+ members producing research to bring rigor, accountability, and standards to the field of AI Evaluations; convened the NeurIPS '24, UC San Diego '25, and ACL '26 workshops on AI evaluation methodology.
 - Hugging Science** 2025 – Now
Co-Founder & Lead (Policy). AI-for-Science initiative; Georgia Channing leads hub curation while I lead the policy track on how scientific AI tools are evaluated, governed, and disclosed across astronomy, biology, climate, and materials science.

New York, NY
- 2024 – Now

University of Connecticut **Academia**
Research Affiliate.
Affiliated with CACC, UCHI, and DMD working on various moonshot projects and advising students.
 - Pluralistic AI Lab (PAL)** 2026 – Now
Co-Founder. Co-founded UConn-affiliated lab researching multi-persona, multi-agent systems that surface diverse and conflicting perspectives, moving beyond single-chatbot interfaces and supporting marginalized communities in democratic deliberation.
 - Reducing Information Ecosystem Threats (RIET) Lab** 2024 – Now
Co-Advisor. Co-advise students with Prof. Shiri Dori-Hacohen on sociotechnical AI alignment.

Storrs, CT

Past Appointments

- 2023 – 24

AdeptID **Industry**
Research Data Scientist.
Audited production hiring ML pipelines for fairness and emerging employment-discrimination compliance (e.g., NYC Local Law 144); developed bias-mitigation strategies for deployed candidate-screening models.

Boston, MA
- 2019 – 23

Northeastern University **Academia**
Doctoral Researcher & Lecturer, Khoury College of Computer Sciences.
Advised by Christo Wilson; promoted from RA to Lecturer (Fall 2023). Designed and taught CS 4973-05: *Responsible Machine Learning* from scratch. Led the pymetrics fairness audit, and contributed to Facebook removing its Special Audiences ad-targeting algorithm.

Boston, MA
- 2021 – 22

Twitter **Industry**
Research Intern, META Team.
On the Machine Ethics, Transparency & Accountability team: studied tradeoffs between demography-agnostic and demography-dependent author-impression fairness; designed home-timeline diversity metrics grounded in user feedback surveys.

San Francisco, CA
- 2020 – 21

Fiddler Labs **Industry**
Research Intern.
Co-developed FairCanary (AIES '22), a threshold-independent intersectional fairness metric using optimal transport, shipped into Fiddler's commercial ML monitoring product; also proposed a unified Shapley framework for explaining ML drift.

Palo Alto, CA
- 2019

LIG, University of Grenoble Alps **Academia**
Visiting Researcher.
Differential information spreading by news organizations on social media; responsibly disclosed an exposed access token to Facebook Bug Bounty.

Grenoble, FR

2014 – 19	IIT Kharagpur Academia	Kharagpur, IN
	Undergraduate Researcher, Complex Networks Research Group.	
	Savitr disaster-response location extraction (WWW '18); news-comment alignment (ECIR '19); legal catchphrase NER.	
2017	Xerox Research Centre Industry	Bangalore, IN
	Research Intern.	
	XTrack BLE vehicle tracking and spatio-temporal surge price prediction; <i>Best Internship Project.</i>	
2016	Google Summer of Code Industry	Remote
	GSoC Student at OpenMRS.	
	Replaced legacy HTML XForms with native REST forms in the OpenMRS Android client; 100K+ lines contributed as top repo contributor.	

PUBLICATIONS

Peer-Reviewed Conference		
ICML '26	Who Evaluates AI's Social Impacts? Mapping Coverage and Gaps in First and Third Party Evaluations	Seoul, KR
	A. Reuel*, Avijit Ghosh* , J. Chim*, A. Tran, Y. Long, J. Mickel, U. Gohar, S. Yadav, P. S. Ammanamanchi, M. Allaham, H. A. Rahmani, M. Akhtar, F. Friedrich, R. Scholz, M. A. Riegler, J. Batzner, E. Habba, A. Saxena, A. Kornilova, K. Wei, P. Soni, Y. Mathew, K. Klyman, J. Sania, S. Sahoo, O. B. Bruvik, P. Sadeghi, S. Goswami, A. Wang, Y. Jernite, Z. Talat, S. Biderman, M. Kochenderfer, S. Koyejo, I. Solaiman	
ICML '26	When AI Benchmarks Plateau: A Systematic Study of Benchmark Saturation	Seoul, KR
	M. Akhtar, A. Reuel, P. Soni, S. Ahuja, P. S. Ammanamanchi, R. Rawal, V. Zouhar, S. Yadav, C. Whitehouse, D. Ki, J. Mickel, L. Choshen, M. Suppa, J. Batzner, J. Chim, J. Sania, Y. Long, H. A. Rahmani, C. Knight, Y. Nan, J. Raj, Y. Fan, S. Singh, S. Sahoo, E. Habba, U. Gohar, S. Pawar, R. Scholz, A. Subramonian, J. Ni, M. Kochenderfer, S. Koyejo, M. Sachan, S. Biderman, Z. Talat, Avijit Ghosh , I. Solaiman	
ICML '26	Position: Agentic Systems Should be General	Seoul, KR
	E. Bandel, A. Yehudai, A. Lacoste, Avijit Ghosh , G. Neubig, M. Mitchell, M. Shmueli-Scheuer, L. Choshen	
FAccT '26	What if AI Systems Weren't Chatbots?	Montréal, CA
	Sourojit Ghosh*, Pranav Narayanan Venkit, Sanjana Gautam, Avijit Ghosh*	
AIES '25	Documenting Patterns of Exoticism of Marginalized Populations within Text-to-Image Generators	Madrid, ES
	Sourojit Ghosh, Sanjana Gautam, Pranav Venkit, Avijit Ghosh	
ICML '25	Stop treating 'AGI' as the north-star goal of AI research	Vancouver, CA
	B. Blili-Hamelin, C. Graziul, L. Hancox-Li, H. Hazan, E. El-Mhamdi, Avijit Ghosh , K. Heller, J. Metcalf, F. Murai, E. Salvaggio, A. Smart, T. Snider, M. Tighanimine, T. Ringer, M. Mitchell, S. Dori-Hacohen	
ICML '25	In-House Evaluation Is Not Enough: Towards Robust Third-Party Flaw Disclosure for General-Purpose AI	Vancouver, CA
	S. Longpre, K. Klyman, R. Appel, S. Kapoor, R. Bommasani, M. Sahar, S. McGregor, Avijit Ghosh , B. Blili-Hamelin, N. Butters, A. Nelson, A. Elazari, A. Sellars, C. Ellis, D. Sherrets, D. Song, H. Geiger, I. Cohen, L. McIlvenny, M. Srikumar, M. Jaycox, M. Anderljung, N. Johnson, N. Carlini, N. Mialhe, N. Marda, P. Henderson, R. Portnoff, R. Weiss, V. Westerhoff, Y. Jernite, R. Chowdhury, P. Liang, A. Narayanan.	
FAccT '25	"It's not a representation of me": Examining Accent Bias and Digital Exclusion in Synthetic AI Voice Services	Athens, GR
	Shira Michel, Sufi Kaur, Sarah Elizabeth Gillespie, Jeffrey Gleason, Christo Wilson, Avijit Ghosh	
AAAI '25	Quantifying Misalignment Between Agents: Towards a Sociotechnical Understanding of Alignment	Philadelphia, US
	Aidan Kierans, Avijit Ghosh , Hananel Hazan, Shiri Dori-Hacohen	
AIES '24	Coordinated Disclosure for AI: Beyond Security Vulnerabilities	San Jose, US
	Sven Cattell, Avijit Ghosh , Lucie-Aimée Kaffee	
WWW '24	Perceptions in pixels: analyzing perceived gender and skin tone in real-world image search results	Singapore, SG
	Jeffrey Gleason, Avijit Ghosh , Christo Wilson	
AIES '23	Bound by the Bounty: Collaboratively Shaping Evaluation Processes for Queer AI Harms	Montréal, CA
	N. Dennler, A. Ovalle, A. Singh, L. Soldaini, A. Subramonian, H. Tu, W. Agnew, Avijit Ghosh , K. Yee, I.F. Peradejordi, Z. Talat, M. Russo, J. Pinhal	
AIES '23	When Fair Classification Meets Noisy Protected Attributes	Montréal, CA
	Avijit Ghosh , Pablo Kvitca, Christo Wilson	
FAccT '23	Queer In AI: A Case Study in Community-Led Participatory AI	Chicago, IL
	Organizers of Queer In AI, A. Ovalle, A. Subramonian, A. Singh, C. Voelcker, D. Sutherland, D. Locatelli, E. Breznik, F. Klubička, H. Yuan, H. J. H. Zhang, J. Shriram, K. Lehman, L. Soldaini, M. Sap, M. Deisenroth, M. Pacheco, M. Ryskina, M. Mundt, M. Agarwal, N. McLean, P. Xu, A. Pranav, R. Korpan, R. Ray, S. Mathew, S. Arora, S. John, T. Anand, V. Agrawal, W. Agnew, Y. Long, Z. Wang, Z. Talat, Avijit Ghosh , N. Dennler, M. Noseworthy, S. Jha, E. Baylor, A. Joshi, N. Bilenko, A. McNamara, R. Gontijo-Lopes, A. Markham, E. Dông, J. Kay, M. Saraswat, N. Vytla, L. Stark.	
FAccT '22	Subverting Fair Image Search with Generative Adversarial Perturbations	Seoul, KR
	Avijit Ghosh , Matthew Jagielski, Christo Wilson	

AIES '22	FairCanary: Rapid Continuous Explainable Fairness <i>Avijit Ghosh*, Aalok Shanbhag*, Christo Wilson</i>	Oxford, UK
AIES '22	Algorithms that “Don’t See Color”: Comparing Biases in Lookalike and Special Ad Audiences <i>Piotr Sapiezynski, Avijit Ghosh, Levi Kaplan, Alan Mislove, Aaron Rieke</i>	Oxford, UK
SIGIR '21	When Fair Ranking Meets Uncertain Inference <i>Avijit Ghosh, Ritam Dutt, Christo Wilson</i>	Montréal, CA
FAccT '21	Building and Auditing Fair Algorithms: A Case Study in Candidate Screening <i>Christo Wilson, Avijit Ghosh, Shan Jiang, Alan Mislove, Lewis Baker, Janelle Szary, Kelly Trindel, Frida Polli</i>	Toronto, CA
ECIR '19	Public Sphere 2.0: Targeted Commenting in Online News Media <i>Ankan Mullick, Sayan Ghosh*, Ritam Dutt*, Avijit Ghosh*, Abhijnan Chakrabarty</i>	Cologne, DE

Peer-Reviewed Journal

Patterns '26	AI for Scientific Discovery is a Social Problem <i>Georgia Channing*, Avijit Ghosh*</i>	Journal
TMLR '26	Open Technical Problems in Open-Weight AI Model Risk Management <i>S. Casper, K. O'Brien, S. Longpre, E. Seger, K. Klyman, R. Bommasani, A. Nrusimha, I. Shumailov, S. Mindermann, S. Basart, F. Rudzicz, K. Pelrine, Avijit Ghosh, A. Strait, R. Kirk, D. Hendrycks, P. Henderson, Z. Kolter, G. Irving, Y. Gal, Y. Bengio, D. Hadfield-Menell</i>	Journal
AAA'21	Connectedness of Markets with Heterogeneous Agents and the Information Cascades <i>Avijit Ghosh, Aditya Chourasiya, Lakshay Bansal, Abhijeet Chandra</i>	Journal

Peer-Reviewed Workshop

T4T '25	Protecting Human Cognition in the Age of AI <i>Anjali Singh, Karan Taneja, Zhitong Guan, Avijit Ghosh</i>	Yokohama, JP
IAAI '25	To Err is AI: A Case Study Informing LLM Flaw Reporting Practices <i>Sean McGregor, Allyson Ettinger, Nick Judd, Paul Albee, Liwei Jiang, Kavel Rao, Will Smith, Shayne Longpre, Avijit Ghosh, Christopher Fiorelli, Michelle Hoang, Sven Cattell, Nouha Dziri</i>	Philadelphia, US
HEGM '22	Can There be Art Without an Artist? <i>Avijit Ghosh, Genoveva Fossas</i>	New Orleans, US
AIDBEI '21	Characterizing Intersectional Group Fairness with Worst-Case Comparisons <i>Avijit Ghosh, Lea Genuit, Mary Reagan</i>	Vancouver, CA
ConPro '19	Analyzing Political Advertisers’ Use of Facebook’s Targeting Features <i>Avijit Ghosh, Giridhari Venkatadri, Alan Mislove</i>	San Francisco, US
SMERP '18	SAVITR: A System for Real-time Location Extraction from Microblogs during Emergencies <i>Ritam Dutt, Kaustubh Hiware, Avijit Ghosh, Rameshwar Bhaskaran</i>	Lyon, FR
WITS '19	WebSelect: A Research Prototype for Optimizing Ad Exposures based on Network Structure <i>Avijit Ghosh, Agam Gupta, Divya Sharma, Uttam Sarkar</i>	Dublin, IE

Book Chapter

Chapter	Evaluating the social impact of generative AI systems in systems and society <i>I. Solaiman, Z. Talat, W. Agnew, L. Ahmad, D. Baker, S. Blodgett, C. Chen, H. Daumé, J. Dodge, I. Duan, E. Evans, F. Friedrich, Avijit Ghosh, U. Gohar, S. Hooker, Y. Jernite, R. Kalluri, A. Lusoli, A. Leidinger, M. Lin, X. Lin, S. Luccioni, J. Mickel, M. Mitchell, J. Newman, A. Ovalle, M. Png, S. Singh, A. Strait, L. Struppek, A. Subramonian.</i>	Oxford Handbook on Generative AI
----------------	--	--

Preprints and Working Manuscripts

Preprint	Economies of Open Intelligence: Tracing Power & Participation in the Model Ecosystem <i>S. Longpre, C. Akiki, C. Lund, A. Kulkarni, E. Chen, I. Solaiman, Avijit Ghosh, Y. Jernite, L. Kaffee</i>
Preprint	Position: Aligning AI Requires Automating Reasoning Norms <i>A. Kierans, Avijit Ghosh, S. Dori-Hacohen</i>
Preprint	Fully Autonomous AI Agents Should Not be Developed <i>Margaret Mitchell, Avijit Ghosh, Alexandra Sasha Luccioni, Giada Pistilli</i>

Preprint

Dual Governance: The intersection of centralized regulation and crowdsourced safety mechanisms for Generative AI

Avijit Ghosh, Dhanya Lakshmi

Preprint

Unified Shapley Framework to Explain Prediction Drift

Aalok Shanbhag, Avijit Ghosh*, Josh Rubin**

Preprint

Supervised extraction of catchphrases from legal documents

Avijit Ghosh, Prerit Gupta*, Ritam Dutt, Kaustubh Hiware, Arpan Mandal, Kripabandhu Ghosh, Saptarshi Ghosh*

** Equal contribution*

AWARDS & GRANTS

Awards

2025	Winner Spotlight Poster	ICML '25
2024	Winner Best AI Art	CVPR '24
2023	Winner Best Paper	FAccT '23
2022	Winner Best Paper - Runner Up	Conpro '22
2019	Winner Best Poster	ECIR '19
2019	Dean's Fellowship for first Year PhD students (\$ 72K)	Northeastern University
2019	Winner Institute Order of Merit - Technology	IIT Kharagpur
2017	Winner Best Internship Project	Xerox Research Centre
2017	Silver Medal Stock Market Analysis	Inter IIT Tech Meet, Kanpur
2016	Gold Medal Software Development	Inter IIT Tech Meet, Mandi
2012	Governor's Medal National Rank 5, ICSE Board	Government of West Bengal
2010	NTSE Scholar National Talent Search Examination	NCERT

Received Grants

2026	SFA Research Grant. (\$9K) — Co-Agents: The Filipino Diaspora Pilot – A Multi-Agent System for Cross-Cultural Democratic Deliberation.	UConn School of Fine Arts
2025	Beyond Polarization: Fostering Plurality in Large Language Model Design and Development. (\$6K)	UCHI AI Seed Fund
2018	SGSIS Institute Challenge Grant. (₹1M)	IIT Kharagpur

Grantmaking & Stewardship

2026	Reviewer for the Sloan Foundation Technology program	Alfred P. Sloan Foundation
2026	Reviewer for AI for Public Goods (AI4PG) Fast Grants	GainForest

TEACHING & MENTORSHIP

Courses Taught

Fall 2023	Responsible Machine Learning (CS 4973-05) <i>Adjunct Professor.</i> Designed and taught this course from scratch — algorithmic fairness, bias auditing, ML privacy, and AI safety, with a real-world fairness term project.	Northeastern University
Spring 2022	Algorithm Audits (CS 4910) <i>Teaching Assistant (Dr. Piotr Sapiezynski).</i> Designing and implementing external audits to investigate harms in deployed ML and AI systems.	Northeastern University

Current Mentees

2026	Afifah Kashif — Visiting Graduate Researcher, Pluralistic AI Lab. MPhil candidate in Human-Inspired AI, University of Cambridge.	University of Connecticut
2026	Isabelle Kohout — Visiting Graduate Researcher, Pluralistic AI Lab.	University of Connecticut
2024	Aidan Kierans — Graduate Researcher, RIET Lab. Ph.D. student in Computer Science & Engineering, co-advised with Prof. Shiri Dori-Hacohen.	University of Connecticut
2025	Kaley Rittichier — Graduate Researcher, RIET Lab. Department of Philosophy.	University of Connecticut

Past Mentees

2024	Ryan Lagasse — Undergraduate Researcher, RIET Lab.	University of Connecticut
2024	Andrew E. Lawrence — Visiting Undergraduate Researcher, RIET Lab. Computer Science, Tufts University.	University of Connecticut

MEDIA MENTIONS

2026	Sarvam AI models excel in Indic, lag frontier systems	Financial Express India
2026	Deutsche Telekom Wants AI Agents to Talk on Your Phone for You	Wired
2025	China leapfrogs US in global market for 'open' AI models	Financial Times
2025	Responsible AI podcast with Alex Alben	The AI Forum
2025	Ethical Machines podcast - "We're Not Ready for Agentic AI"	Ethical Machines
2025	Can AI take your coding job?	NPR Marketplace
2025	California finally beats Big Tech in court	Politico
2025	Hugging-face experts warn: We shouldn't give AI agents full control	T3N (German)
2025	Why handing over total control to AI agents would be a huge mistake	MIT Technology Review Op-Ed
2025	The hidden cost of brainstorming with ChatGPT	Business Insider
2025	DeepSeek pioneers a new way for AI to 'reason'	Science News Explores
2025	Two founders built a jobs board for AI agents. Humans need not apply — but their skills are still required.	Business Insider
2025	Why truly open-source AI remains out of reach	HT TechCircle
2025	What does OpenAI get from Stargate? A \$500 billion chance to build a whole new moat.	Business Insider
2025	OpenAI's Stargate may be tech's biggest gamble ever, but here's what's really at stake	Fortune
2025	AI will transform everything from daily life to businesses	Dainik Bhaskar Op-ed (Hindi)
2024	This Wearable AI Notetaker Will Transcribe Your Meetings — and Someday, Your Entire Life	Wired
2024	In California, the controversial AI protection law is about to be passed	Les Echos (French)
2024	We finally have a definition for open-source AI	MIT Technology Review
2024	World's biggest hacker fest spotlights AI's soaring importance in the high-stakes cybersecurity war—and its vulnerability	Fortune
2024	AI full of prejudices both a challenge and an opportunity for India	Dainik Bhaskar Op-ed (Hindi)
2024	Common image search results are overwhelmingly white, a new study finds	Fast Company
2024	As AI tools get smarter, they're growing more covertly racist, experts find	The Guardian
2024	LLMs become more covertly racist with human intervention	MIT Technology Review
2023	How Can AI Affect LGBTQIA+ People In India? Three Indian-Origin Queer Researchers Explain	IndiaTimes
2023	He hacked AI chatbots to find flaws and vulnerabilities. Now Northeastern's Avijit Ghosh is writing a report on combating these problems	Northeastern Global News
2023	Radio Interview - AirTalk	LAIST/NPR
2023	Radio Interview - As It Happens	CBC Canada
2023	When Hackers Descended to Test A.I., They Found Flaws Aplenty	The New York Times
2023	From accessibility efforts to ethical concerns, here are our AI takeaways from SXSW content creators should consider	Passionfruit
2021	NYC aims to be first to rein in AI hiring tools	Associated Press
2021	Auditors are testing hiring algorithms for bias, but there's no easy fix	MIT Technology Review
2021	New York City Proposes Regulating Algorithms Used in Hiring	Wired
2021	Supporting Responsible Use of AI and Equitable Outcomes in Financial Services	The Federal Reserve
2019	Facebook Ads Can Still Discriminate Against Women and Older Workers, Despite a Civil Rights Settlement	Propublica
2019	Facebook Agreed Not to Let Its Ads Discriminate. But They Still Can.	Mother Jones

SPEAKING ENGAGEMENTS

2026	Evaluation Reporting in the Age of Agentic AI	Agentic Evaluations Workshop, Hugging Face
2026	What if AI Systems Weren't Chatbots?	Weizenbaum Institute, Berlin
2026	Opening panel and fireside chat at IASEAI 2026	IASEAI, Paris
2025	Panel: "AI and the Creative Industry"	Canadian Consulate and All Tech is Human
2025	Panel: "Rights (to Life, Liberty, and the Pursuit of Happiness)"	Uconn Humanities Institute AI Symposium
2025	Panel: "Is U.S. Policy Ready for Agentic AI?"	Center for Data Innovation
2025	Panel on Open Source AI Regulation at the Summit on State AI Legislation	Digital Ethics Center, Yale University
2025	NLP and Generative AI Panel	Harvard Business School Tech Conference
2025	Coordinated Disclosure for AI: Beyond Security Vulnerabilities	The MIT Sloan AI Conference
2024	Coordinated Disclosure for AI: Beyond Security Vulnerabilities	Future of Third-Party AI Evaluations
2024	Bridging the Gap: Real-World AI Biases and Responsible Governance Frameworks	IIT Kharagpur

2024	ERASED BY AI: Personal Experiences of the Psychological Impact of AI Bias	Arthur AI Fest
2024	Teaching Responsible AI: Building with open source and Hugging Face	Northeastern University
2024	Coordinated Disclosure for AI: Beyond Security Vulnerabilities	EQUAL lab at MILA
2024	Technology Impact on Cybersecurity	Boston University
2024	AI Vulnerability Reporting Event in the US Congress	Hackers on the Hill
2023	Queer Bias Bounty: Considerations for community audits and lessons learnt	Centre for Data Ethics and Innovation, UK
2023	Can There be AI Art Without an Artist?	South By Southwest (SXSW)
2023	On the evolving tension between centralized regulation and decentralized development of Text-to-Image Models	AIDBEI at AAAI
2022	Proxies for bias monitoring: Ethics workshop	Centre for Data Ethics and Innovation, UK
2022	Subverting Fair Image Search with Generative Adversarial Perturbations as part of the 'Celebrating Young Researchers' event	Trustworthy ML Initiative

LEADERSHIP AND SERVICE

2026	Pluralistic AI Lab (PAL), University of Connecticut	Co-Founder (with Dr. Zhenzhen Qi)
2025	Hugging Science	Co-Founder & Lead (Policy)
2024	EvalEval Coalition	Founder & Lead
2026	ACL Workshop on Evaluating AI in Practice	General Chair
2026	International Conference on Machine Learning (ICML)	Program Committee
2026	ACM Conference on Fairness, Accountability, and Transparency (FAccT)	Program Committee
2026	AAAI Conference on Artificial Intelligence	Program Committee
2025	Workshop on Evaluating AI in Practice, UC San Diego	General Chair
2025	ACM Conference on Fairness, Accountability, and Transparency	Registration Chair
2025	AAAI/ACM Conference on AI, Ethics, and Society	Program Committee
2025	ACM Conference on Fairness, Accountability, and Transparency	Program Committee
2025	Conference on Language Modeling	Program Committee
2024	EvalEval: Evaluating Evaluations: Examining Best Practices for Measuring Broader Impacts of Generative AI	General Chair
2024	AAAI/ACM Conference on AI, Ethics, and Society	Program Committee
2024	Conference on Neural Information Processing Systems	Program Committee
2024	ICWSM: The International AAAI Conference on Web and Social Media	Program Committee
2024	The Web Conference	Program Committee
2024	ACM Conference on Fairness, Accountability, and Transparency	Program Committee
2023	SIGIR DEI lunch with speaker panel on disability in computing	Chair
2023	FAccT CRAFT Workshop on an India-first Responsible AI research agenda	Organizer
2023	FAccTRec: Workshop on Responsible Recommendation	Program Committee
2023	Conference on Neural Information Processing Systems	Program Committee
2023	AAAI/ACM Conference on AI, Ethics, and Society	Program Committee
2023	ACM Conference on Fairness, Accountability, and Transparency	Program Committee
2023	The Web Conference	Program Committee
2022	FAccT CRAFT Workshop on Humanitarian AI for the Global South	Organizer
2022	FAccT CRAFT Workshop on Identifying Queer Harms as a bias bounty with Queer in AI	Organizer
2022	ACM Conference on Fairness, Accountability, and Transparency	Program Committee
2022	AAAI/ACM Conference on AI, Ethics, and Society	Program Committee
2022	Conference on Neural Information Processing Systems	Program Committee
2022	Conference on Empirical Methods in Natural Language Processing	Program Committee
2021	SIGIR Queer in AI social with speaker panel on queer stereotypes in web search	Organizer
2021	Conference on Neural Information Processing Systems	Program Committee
2017	Kharagpur Winter Of Code	Founder
2016	Kharagpur Open Source Society	Founder