

Avijit Ghosh

Applied Policy Researcher, ML & Society, Hugging Face

☎ (+1) 857-337-0180 | ✉ avijit@huggingface.co (Work) | avijitg22@gmail.com (Personal) | 🌐 evijit.io | in [evijit](#) | 📺 [evijit](#)

Algorithmic Fairness Ethical AI Machine Learning AI Explainability Policy Computational Social Science

Education

Northeastern University

Ph.D. in Computer Science (Advised by Dr. Christo Wilson)

Boston, MA

2019 - 2023

Indian Institute of Technology (IIT) Kharagpur

B.Tech. in Chemical Engineering, M.Tech in Financial Engineering, Minor in Computer Science

Kharagpur, India

2014 - 2019

Doctoral Thesis

Algorithmic Fairness in the Real World: Challenges and Considerations

Defended June 2023

- Social bias in machine learning algorithms is a widespread problem that has been addressed through various measures, but implementing fair machine learning systems in the real world is challenging due to issues like noisy demographic information, adversarial vulnerabilities, policy restrictions and complex interactions between humans and algorithms.
- In my thesis, I attempt to outline these problems in fair ML systems, with the aim to gain a more complete understanding of the challenges involved and to be able to provide technical and policy recommendations to overcome their real world implementation challenges.

Publications

Peer-Reviewed Conference

"It's not a representation of me": Examining Accent Bias and Digital Exclusion in Synthetic AI Voice Services

FACCT'25

Shira Michel, Sufi Kaur, Sarah Elizabeth Gillespie, Jeffrey Gleason, Christo Wilson, **Avijit Ghosh**

Athens, Greece

Quantifying Misalignment Between Agents: Towards a Sociotechnical Understanding of Alignment

AAAI '25

Aidan Kierans, **Avijit Ghosh**, Hananel Hazan, Shiri Dori-Hacohen

Philadelphia, USA

Coordinated Disclosure for AI: Beyond Security Vulnerabilities

AIES '24

Sven Cattell, **Avijit Ghosh**, Lucie-Aimée Kaffee

San Jose, USA

Perceptions in pixels: analyzing perceived gender and skin tone in real-world image search results

WWW '24

Jeffrey Gleason, **Avijit Ghosh**, Christo Wilson

Singapore

Bound by the Bounty: Collaboratively Shaping Evaluation Processes for Queer AI Harms

AIES '23

N. Dennler, A. Ovalle, A. Singh, L. Soldaini, A. Subramonian, H. Tu, W. Agnew, **Avijit Ghosh**, K. Yee,

Montreal, Canada

I.F. Peradejordi, Z. Talat, M. Russo, J. Pinhal

When Fair Classification Meets Noisy Protected Attributes

AIES '23

Avijit Ghosh, Pablo Kvitca, Christo Wilson

Montreal, Canada

Queer In AI: A Case Study in Community-Led Participatory AI

FACCT '23

Organizers of Queer In AI (50 authors)

Chicago, Illinois

Subverting Fair Image Search with Generative Adversarial Perturbations

FACCT '22

Avijit Ghosh, Matthew Jagielski, Christo Wilson

Seoul, South Korea

FairCanary: Rapid Continuous Explainable Fairness

AIES '22

Avijit Ghosh*, Aalok Shanbhag*, Christo Wilson

Oxford, United Kingdom

Algorithms that "Don't See Color": Comparing Biases in Lookalike and Special Ad Audiences

AIES '22

Piotr Sapiezynski, **Avijit Ghosh**, Levi Kaplan, Alan Mislove, Aaron Rieke

Oxford, United Kingdom

When Fair Ranking Meets Uncertain Inference

SIGIR '21

Avijit Ghosh, Ritam Dutt, Christo Wilson

Montreal, Canada / Virtual

Building and Auditing Fair Algorithms: A Case Study in Candidate Screening

FACCT '21

Christo Wilson, **Avijit Ghosh**, Shan Jiang, Alan Mislove, Lewis Baker, Janelle Szary, Kelly Trindel, Frida Polli

Toronto, Canada / Virtual

Public Sphere 2.0: Targeted Commenting in Online News Media

ECIR '19

Ankan Mullick, Sayan Ghosh*, Ritam Dutt*, **Avijit Ghosh***, Abhijnan Chakrabarty

Cologne, Germany

Peer-Reviewed Workshop

Protecting Human Cognition in the Age of AI

Anjali Singh, Karan Taneja, Zhitong Guan, **Avijit Ghosh**

To Err is AI: A Case Study Informing LLM Flaw Reporting Practices

Sean McGregor, Allyson Ettinger, Nick Judd, Paul Albee, Liwei Jiang, Kavel Rao, Will Smith, Shayne Longpre, **Avijit Ghosh**, Christopher Fiorelli, Michelle Hoang, Sven Cattell, Nouha Dziri

Can There be Art Without an Artist?

Avijit Ghosh, Genoveva Fossas

Characterizing Intersectional Group Fairness with Worst-Case Comparisons

Avijit Ghosh, Lea Genuit, Mary Reagan

Analyzing Political Advertisers' Use of Facebook's Targeting Features

Avijit Ghosh, Giridhari Venkatadri, Alan Mislove

SAVITR: A System for Real-time Location Extraction from Microblogs during Emergencies

Ritam Dutt, Kaustubh Hiware, **Avijit Ghosh**, Rameshwar Bhaskaran

WebSelect: A Research Prototype for Optimizing Ad Exposures based on Network Structure

Avijit Ghosh, Agam Gupta, Divya Sharma, Uttam Sarkar

Tools4Thoughts@CHI '25

Yokohama, Japan

IAAI@AAAI '25

Philadelphia, USA

HEGM@NeurIPS '22

New Orleans, USA

AIDBEI@AAAI '21

Vancouver, Canada / Virtual

Conpro@S&P '19

San Francisco, USA

SMERP@WWW '18

Lyon, France

WITS'19

Dublin, Ireland

Peer-Reviewed Journal

Connectedness of Markets with Heterogeneous Agents and the Information Cascades

Avijit Ghosh, Aditya Chourasiya, Lakshay Bansal, Abhijeet Chandra

AAA'21

Journal

Book Chapter

Evaluating the social impact of generative AI systems in systems and society

Irene Solaiman, Zeerak Talat, et al. (including **Avijit Ghosh**)

Chapter

Oxford Handbook on Generative AI (forthcoming)

Preprints and Working Manuscripts

Fully Autonomous AI Agents Should Not be Developed

Margaret Mitchell, **Avijit Ghosh**, Alexandra Sasha Luccioni, Giada Pistilli,

Preprint

Stop treating 'AGI' as the north-star goal of AI research

B. Blili-Hamelin, C. Graziul, L. Hancox-Li, H. Hazan, E. El-Mhamdi, **Avijit Ghosh**, K. Heller, J. Metcalf, F. Murai, E. Salvaggio, A. Smart, T. Snider, M. Tighanimine, T. Ringer, M. Mitchell, S. Dori-Hacohen

Preprint

In-House Evaluation Is Not Enough: Towards Robust Third-Party Flaw Disclosure for General-Purpose AI

Shayne Longpre, Kevin Klyman, Ruth E. Appel, Sayash Kapoor, Rishi Bommasani, Michelle Sahar, Sean McGregor, **Avijit Ghosh**, Borhane Blili-Hamelin, Nathan Butters, Alondra Nelson, Amit Elazari, Andrew Sellars, Casey John Ellis, Dane Sherrets, Dawn Song, Harley Geiger, Ilona Cohen, Lauren McIlvenny, Madhulika Srikumar, Mark M. Jaycox, Markus Anderljung, Nadine Farid Johnson, Nicholas Carlini, Nicolas Miallhe, Nik Marda, Peter Henderson, Rebecca S. Portnoff, Rebecca Weiss, Victoria Westerhoff, Yacine Jernite, Rumman Chowdhury, Percy Liang, Arvind Narayanan

Preprint

Dual Governance: The intersection of centralized regulation and crowdsourced safety mechanisms for Generative AI

Avijit Ghosh, Dhanya Lakshmi

Preprint

Unified Shapley Framework to Explain Prediction Drift

Aalok Shanbhag*, **Avijit Ghosh***, Josh Rubin*

Preprint

Supervised extraction of catchphrases from legal documents

Avijit Ghosh*, Prerit Gupta*, Ritam Dutt, Kaustubh Hiware, Arpan Mandal, Kripabandhu Ghosh, Saptarshi Ghosh

Term paper

* Equal contribution

Grants

2025	Beyond Polarization: Fostering Plurality in Large Language Model Design and Development. (\$6K)	UCHI AI Seed Fund
2018	SGSIS Institute Challenge Grant. (₹1M)	IIT Kharagpur

Awards

2024	Winner Best AI Art	CVPR'24
2023	Winner Best Paper	FAccT'23
2022	Winner Best Paper - Runner Up	Conpro'22
2019	Winner Best Poster Award	ECIR'19
2019	Dean's Fellowship for first Year PhD students (USD 72K)	Northeastern University
2019	Winner Institute Order of Merit - Technology	IIT Kharagpur
2017	Silver Medal Stock Market Analysis	Inter IIT Tech Meet, Kanpur
2016	Gold Medal Software Development	Inter IIT Tech Meet, Mandi
2012	Governor's Medal National Rank 5, ICSE Board	Government of West Bengal
2010	NTSE Scholar National Talent Search Examination	NCERT

Teaching

Responsible Machine Learning

Northeastern University

Lecturer

Fall 2023

- Explores the ethical challenges and responsibilities of creating and deploying machine learning (ML) models.
- Biases in ML models, methods for uncovering them, and algorithmic fairness techniques to mitigate them
- Term project to apply algorithmic fairness to a real world scenario

Algorithmic Auditing

Northeastern University

Teaching Assistant for Dr. Piotr Sapiezynski

Spring 2022

- Designing audits that measure the effects of interests and control noise sources
- Minimize potential harms of audits to all stakeholders
- Legal bounds of algorithm audits
- Beyond audits: potential harms that cannot be measured through audits

Academic Experience

University of Connecticut

Storrs, CT

Associate Researcher at Connecticut Advanced Computing Center

Mar 2024 – Present

- Collaborating with Prof. Shiri Dori-Hacohen in the Risk and Information Ecosystem Threats (RIET) Lab on research and mentoring PhD students and Undergraduate Researchers on AI Alignment related projects as a visiting/associate expert.

Northeastern University

Boston, MA

Lecturer at Khoury College of Computer Sciences

Sep 2023 – Present

- Teaching CS 4973-05 Responsible Machine Learning to senior undergrads. With the help of readings, guest lectures and original course material, the focus of the course is to empower students to responsibly deploy ethical and fair machine learning models for societal benefit.

Northeastern University

Boston, MA

Research Assistant at Khoury College of Computer Sciences

Sep 2019 – Present

- Analyzing Fair ranking systems and showing how they fail in the presence of noisy protected attribute data. Also investigated adversarial attacks on the guaranteed fairness of retrieval systems.
- A cooperative fairness audit of the recommendation algorithm of [PyMetrics](#), a talent matching software. [Press Release](#).
- Investigated Facebook's Special Audiences system for opportunity advertisements and showed that the audience creation algorithm was still biased against women, seniors and minorities.
- Analyzed the ad reach and spend information obtained from Facebook's ad transparency feature and the personal targeting dataset from Propublica's Facebook ad dataset and showed that advertisers with higher budgets use more privacy sensitive targeting techniques like PII or Lookalike audiences. Findings published and presented at [IEEE ConPro 2019](#).

LIG, University of Grenoble Alps

Grenoble, France

Visiting Researcher

May 2019 – July 2019

- Study of how news companies promote different items on social media, investigating possible patterns of differential information spreading using both posts and ads.
- We also discovered and reported an exposed access token bug to [Facebook Bug Bounty](#).

IIT Kharagpur

Kharagpur, India

Undergraduate Researcher - Complex Networks Research Group

2014 – 2019

- Automated Extraction of Catchwords from Legal Documents using a novel NER tagger to help categorize lengthy legal texts.
- Automatically position user comments against relevant news article paragraphs. Presented at [ECIR 2019](#).
- Savitr - A real-time location extraction system for disaster management using twitter. Presented at [WWW-SMERP 2018](#).
- Classification and Summarization of tweets during a disaster event, presented at [IBM Day 2016](#).

Industry Experience

Hugging Face

New York, NY

Applied Policy Researcher

Mar 2024 – Present

- The Applied Policy Researcher position at Hugging Face involves working within the Machine Learning and Society team to bridge the gap between regulatory and technical realms. The role centers on developing tools to audit ML biases and engaging in policy discussions to facilitate understanding between policymakers and developers, with a focus on democratizing access to advanced machine learning technology.
- Responsibilities include contributing to ongoing public governance efforts by evaluating the social impacts of technology, providing feedback on regulatory proposals, and collaborating on projects such as governance in the BigCode project and evaluating the social impact of generative AI systems.

AdeptID

Research Data Scientist

Boston, MA

Jul 2023 – Feb 2024

- Work with the Data Science and Engineering teams and external policy experts to audit internal systems and ensure that AdeptID's ML models are fair and unbiased and adhere to regulations.
- Conduct original research on ML Fairness and investigate solutions to unanswered questions about potential sources of bias in an industrially deployed ML pipeline.

Twitter

Research Intern

San Francisco, CA

Sep – Dec 2021 and Jun – Aug 2022

- Worked with the META (Machine Ethics, Transparency and Accountability) team at Twitter, to investigate the relationship between demography agnostic and demography dependent author impression fairness metrics at scale.
- Developed home timeline diversity metrics based on user feedback, to find balance between recommendation efficiency and fairness.

Fiddler Labs

Research Intern

Palo Alto, CA

Oct 2020 – Apr 2021

- Explain distributional shifts in Machine Learning model outputs by unifying Shapley based methods.
- Using optimal transport theory, proposed a threshold independent fairness metric that allows for real time explanations.
- Worked with the product team and civil rights lawyers in the deployment of Fiddler's Machine Learning model fairness dashboard. Introduced and incorporated intersectional fairness metrics in the product.

Xerox Research Centre

Research Intern

Bangalore, India

May 2017 – July 2017

- Implemented XTrack, a Smart Vehicle Tracking and Battery usage minimizing Algorithm, using BLE to relay GPS information.
- Proposed a method for Uber-like Surge Price Prediction using Spatio-Temporal techniques like the Neural Hawkes and Recurrent Marked Temporal Point Process. Awarded the title of [Best Internship Project](#).

Google Summer of Code

GSoC Student at OpenMRS

Remote

Apr 2016 - Aug 2016

- Replaced the HTML XForms system used with native generated forms using the Forms REST Api in the android client of the Opensource Medical Record System. Added offline form saving. Configured Travis CI to automatically build and push the apk to the play store.
- Overall, contributed [100K lines of code](#) and became the top code contributor in the project repository.

Academic Service

2025	ACM Conference on Fairness, Accountability, and Transparency	Registration Chair
2025	Conference on Language Modeling	Program Committee
2024	EvalEval: Evaluating Evaluations: Examining Best Practices for Measuring Broader Impacts of Generative AI	General Chair
2024	Conference on Neural Information Processing Systems	Program Committee
2024	ICWSM: The International AAAI Conference on Web and Social Media	Program Committee
2024	The Web Conference	Program Committee
2024	ACM Conference on Fairness, Accountability, and Transparency	Program Committee
2023	FAccTRec: Workshop on Responsible Recommendation	Program Committee
2023	Conference on Neural Information Processing Systems	Program Committee
2023	AAAI/ACM Conference on AI, Ethics, and Society	Program Committee
2023	ACM Conference on Fairness, Accountability, and Transparency	Program Committee
2023	The Web Conference	Program Committee
2022	ACM Conference on Fairness, Accountability, and Transparency	Program Committee
2022	AAAI/ACM Conference on AI, Ethics, and Society	Program Committee
2022	Conference on Neural Information Processing Systems	Program Committee
2022	Conference on Empirical Methods in Natural Language Processing	Program Committee
2021	Conference on Neural Information Processing Systems	Program Committee

Outreach and Leadership

2023	SIGIR DEI lunch with speaker panel on disability in computing	Chair
2023	FAccT CRAFT Workshop on an India-first Responsible AI research agenda	Organizer
2022	FAccT CRAFT Workshop on Humanitarian AI for the Global South	Organizer
2022	FAccT CRAFT Workshop on Identifying Queer Harms as a bias bounty with Queer in AI	Organizer
2021	SIGIR Queer in AI social with speaker panel on queer stereotypes in web search	Organizer
2017	Kharagpur Winter Of Code	Founder
2016	Kharagpur Open Source Society	Founder

Speaking Engagements

2025	Panel: "Is U.S. Policy Ready for Agentic AI?"	Center for Data Innovation
2025	Panel on Open Source AI Regulation at the Summit on State AI Legislation	Digital Ethics Center, Yale University
2025	NLP and Generative AI Panel	HBS Tech Conference
2025	Coordinated Disclosure for AI: Beyond Security Vulnerabilities	The MIT Sloan AI Conference
2024	Coordinated Disclosure for AI: Beyond Security Vulnerabilities	The Future of Third-Party AI Evaluation Workshop
2024	Bridging the Gap: Real-World AI Biases and Responsible Governance Frameworks	IIT Kharagpur
2024	ERASED BY AI: Personal Experiences of the Psychological Impact of AI Bias	Arthur AI Fest
2024	Teaching Responsible AI: Building with open source and Hugging Face	Northeastern University
2024	Coordinated Disclosure for AI: Beyond Security Vulnerabilities	EQUAL lab at MILA
2024	Technology Impact on Cybersecurity	Boston University
2024	AI Vulnerability Reporting Event in the US Congress	Hackers on the Hill
2023	Queer Bias Bounty: Considerations for community audits and lessons learnt	Centre for Data Ethics and Innovation, UK
2023	Can There be AI Art Without an Artist?	South By Southwest (SXSW)
2023	On the evolving tension between centralized regulation and decentralized development of Text-to-Image Models	AIDBEI at AAAI
2022	Proxies for bias monitoring: Ethics workshop	Centre for Data Ethics and Innovation, UK
2022	Subverting Fair Image Search with Generative Adversarial Perturbations as part of the 'Celebrating Young Researchers' event	Trustworthy ML Initiative

Media Mentions

2025	California finally beats Big Tech in court	Politico
2025	Hugging-face experts warn: We shouldn't give AI agents full control	T3N Germany
2025	Why handing over total control to AI agents would be a huge mistake	MIT Technology Review Op-Ed
2025	The hidden cost of brainstorming with ChatGPT	Business Insider
2025	DeepSeek pioneers a new way for AI to 'reason'	Science News Explores
2025	Two founders built a jobs board for AI agents. Humans need not apply — but their skills are still required.	Business Insider
2025	Why truly open-source AI remains out of reach	HT TechCircle
2025	What does OpenAI get from Stargate? A \$500 billion chance to build a whole new moat.	Business Insider
2025	OpenAI's Stargate may be tech's biggest gamble ever, but here's what's really at stake	Fortune
2025	AI will transform everything from daily life to businesses	Dainik Bhaskar Op-ed
2024	This Wearable AI Notetaker Will Transcribe Your Meetings — and Someday, Your Entire Life	Wired
2024	En Californie, la loi de protection contre l'IA amendée pour apaiser la Silicon Valley	Les Echos
2024	We finally have a definition for open-source AI	MIT Technology Review
2024	World's biggest hacker fest spotlights AI's soaring importance in the high-stakes cybersecurity war—and its vulnerability	Fortune
2024	AI full of prejudices both a challenge and an opportunity for India	Dainik Bhaskar Op-ed
2024	Common image search results are overwhelmingly white, a new study finds	Fast Company
2024	As AI tools get smarter, they're growing more covertly racist, experts find	The Guardian
2024	LLMs become more covertly racist with human intervention	MIT Technology Review
2023	How Can AI Affect LGBTQIA+ People In India? Three Indian-Origin Queer Researchers Explain	IndiaTimes
2023	He hacked AI chatbots to find flaws and vulnerabilities. Now Northeastern's Avijit Ghosh is writing a report on combating these problems	Northeastern Global News
2023	Radio Interview - AirTalk	LAIST/NPR
2023	Radio Interview - As It Happens	CBC Canada
2023	When Hackers Descended to Test A.I., They Found Flaws Aplenty	The New York Times
2023	From accessibility efforts to ethical concerns, here are our AI takeaways from SXSW content creators should consider	Passionfruit
2021	NYC aims to be first to rein in AI hiring tools	Associated Press
2021	Auditors are testing hiring algorithms for bias, but there's no easy fix	MIT Technology Review
2021	New York City Proposes Regulating Algorithms Used in Hiring	Wired
2021	Supporting Responsible Use of AI and Equitable Outcomes in Financial Services	The Federal Reserve

2019 **Facebook Ads Can Still Discriminate Against Women and Older Workers, Despite a Civil Rights Settlement**

ProPublica

2019 **Facebook Agreed Not to Let Its Ads Discriminate. But They Still Can.**

Mother Jones